

# Building Containers for Python Applications

Moshe Zadka – <https://cobordism.com>

2021

# Acknowledgement of Country

Belmont (in San Francisco Bay Area Peninsula)  
Ancestral homeland of the Ramaytush Ohlone people

# What is good

# What is good

- ▶ To crush your enemies

# What is good

- ▶ To crush your enemies
- ▶ To see them driven before you

# What is good

- ▶ To crush your enemies
- ▶ To see them driven before you
- ▶ Um, wrong slides

# What is good

► Fast

# What is good

- ▶ Fast
- ▶ Small



# What is good

- ▶ Fast
- ▶ Small
- ▶ Secure

# What is good

- ▶ Fast
- ▶ Small
- ▶ Secure
- ▶ Usable

# Specifying the requirements

Let's be more concrete

- ▶ Keep up to date

# Specifying the requirements

Let's be more concrete

- ▶ Keep up to date
- ▶ Reproducible builds

# Specifying the requirements

Let's be more concrete

- ▶ Keep up to date
- ▶ Reproducible builds
- ▶ No compilers in prod

# Specifying the requirements

Let's be more concrete

- ▶ Keep up to date
- ▶ Reproducible builds
- ▶ No compilers in prod
- ▶ Keep size (reasonably) small

# Up to date

- ▶ Install security updates

# Up to date

- ▶ Install security updates
- ▶ But when?



# Reproducible builds

Same code gives same results

# Reproducible builds

Same code gives same results  
...mostly

# No compilers in prod

A common anti-pattern

# No compilers in prod

A common anti-pattern  
...surprisingly easy to get wrong!

# Size

- ▶ Diminishing returns

# Size

- ▶ Diminishing returns
- ▶ Cost savings

# Support binary wheels

Installing and building

# Support binary wheels

Installing and building  
Faster



# Support binary wheels

Installing and building  
Faster  
Simplifies images

# Not run as root

## General hygiene

# Minimal privileges

Especially avoid permissions to `pip install`

# Fast rebuilds

Responsiveness!

# Base OS

The distro wars are back?

# Base - size

Most modern distros have a decent minimal server

# Base - size

Most modern distros have a decent minimal server  
...but Debian is easiest to get smallest.

# Base - LTS/support

Usually around 5 years



# Base - LTS/support

Usually around 5 years

Gives you time to upgrade!

# Base - Volatility

How much change?

Security? Backports? Fixes?

# Debian

LTS: 5 years  
Conservative

LTS: 5 years  
(Universe, Multiverse, etc...)  
Fairly conservative

# Alpine (probably not)

Uses musl, not manylinux compatible

# Rolling releases (probably not)

Up to date, but...

# Rolling releases (probably not)

Up to date, but...  
updates can change major versions!

Rolling release!



# How to get Python?

So many options...

# Not system Python

Distros aim Python at distro packages

# Not system Python

Distros aim Python at distro packages  
not user programs.

# Appropriate repositories

Famous examples: deadsnakes PPA for Ubuntu

Builds and installs Python

# python-build

Builds and installs Python

# Source

```
RUN configure [...]
```

```
RUN make
```

```
RUN make install
```

# Trade-offs

Control vs. Work vs. Problems



# Versions

Support multiple for upgrade path

# Versions

Support multiple for upgrade path  
2-3

# Docker multistage (quick recap)

Only one stage output

# Docker multistage (quick recap)

Only one stage output  
other stages help

# FROM

Use previous stage as starting image

# COPY -from

Copy files from previous stage

## Stages as modules

```
FROM ubuntu as security-updates
RUN add-apt-repository ppa:deadsnakes/ppa
RUN apt-get update
RUN apt-get upgrade
```

```
FROM security-updates as with-38
RUN apt-get install python3.8
```

```
FROM security-updates as with-39
RUN apt-get install python3.9
```

# Separate build and runtime

Especially when building from source!



# Separate build and runtime

Especially when building from source!

```
FROM ubuntu as builder
# install build dependencies
# build Python into /opt/myorg/python
```

```
FROM ubuntu as runtime
COPY —from=builder \
    /opt/myorg/python \
    /opt/myorg/python
```

# Optimizing layers

Put everything under /opt/myorg  
Use one COPY --from=...

# Optimizing size

After building Python, remove:

- ▶ Tests
- ▶ Builder dependencies (in runtime)
- ▶ ....and more

# Binary wheels

- ▶ Build with builder
- ▶ Copy to runtime
- ▶ Install in virtual environment

# Binary wheels (alt)

- ▶ Build with builder
- ▶ Install in virtual environment
- ▶ Copy virtual environment to runtime

# Patchelf

Used to make wheels self-contained  
Newst version needed

# Auditwheel

Use pip to install

# Self-contained binary wheels

Run

```
auditwheel repair —platform linux_x86_64
```



# Self-contained binary wheels

Run

```
auditwheel repair —platform linux_x86_64
```

No need for binary dependencies!

# Portable binary wheels

- ▶ Oldest supported?

# Portable binary wheels

- ▶ Oldest supported?

Example:

```
auditwheel repair --platform manylinux_2_27_x86_64
```

# Generating binary wheels

Build instructions in docs

# Generating binary wheels

Build instructions in docs

Build dependencies

# Optimizing layers

Reduce copies

# Optimizing layers

Reduce copies

Prep

# Optimizing caching

Where to build wheel?



# Optimizing caching

Where to build wheel?

What invalidates caching?

# Conclusion

- ▶ Wrong easier than right

# Conclusion

- ▶ Wrong easier than right
- ▶ But right is amazing

# Conclusion

- ▶ Wrong easier than right
- ▶ But right is amazing
- ▶ Think before you docker

## Further Resources

Itamar's series – <https://pythonspeed.com/docker/>